











A Tutorial on Wikimedia Visual Resources and its Application to Neural Visual Recommender Systems

Denis Parra¹, Antonio Ossa-Guerra¹, Manuel Cartagena¹, **Patricio Cerda-Mardini**², Felipe del Río¹, Isidora Palma¹, Diego Saez-Trumper³, and Miriam Redi³

- 1. Pontificia Universidad Católica de Chile
- 2. mindsdb
- 3. Wikimedia Foundation













Session 3: Visual BPR methods

Denis Parra¹, Antonio Ossa-Guerra¹, Manuel Cartagena¹, **Patricio Cerda-Mardini**², Felipe del Río¹, Isidora Palma¹, Diego Saez-Trumper³, and Miriam Redi³

- 1. Pontificia Universidad Católica de Chile
- 2. mindsdb
- 3. Wikimedia Foundation



VBPR: Visual Bayesian Personalized Ranking from Implicit Feedback

Ruining He UC San Diego r4he@ucsd.edu Julian McAuley UC San Diego jmcauley@ucsd.edu

Context

- Previous work did not consider the *visual appearance* of the items

- Modeling derived from BPR, suitable to uncover visual factors

Addressing cold-start problem

VBPR Key insights

- Each item is represented by its visual features, from a pretrained AlexNet

- Preference predictor is a complex term that can be simplified thanks to BPR

Model: Preference predictor

$$\widehat{x}_{u,i} = \alpha + \beta_u + \beta_i + \gamma_u^T \gamma_i + \theta_u^T (\mathbf{E} f_i) + \beta'^T f_i$$

$$lpha$$
 Global offset
$$eta_u + eta_i$$
 Bias terms
$$\gamma_u^T \gamma_i$$
 Latent factors (compatibility between user and item)

Model: Preference predictor

$$\widehat{x}_{u,i} = \alpha + \beta_u + \beta_i + \gamma_u^T \gamma_i + \theta_u^T (\mathbf{E} f_i) + \beta'^T f_i$$

$$heta_u^T(\mathbf{E}f_i)$$
 (user preference over visual dimensions)

 $\beta'^T f_i$

Users' visual bias
(user's overall opinion towards visual appearance of a given item)

Model: Preference predictor while training

$$\widehat{x}_{uij} = \widehat{x}_{u,i} - \widehat{x}_{u,j}$$

$$\widehat{x}_{u,i} = \alpha + \beta_u + \beta_i + \gamma_u^T \gamma_i + \theta_u^T (\mathbf{E} f_i) + \beta'^T f_i$$

$$\widehat{x}_{u,j} = \alpha + \beta_u + \beta_j + \gamma_u^T \gamma_j + \theta_u^T (\mathbf{E} f_j) + \beta'^T f_j$$

Model diagram

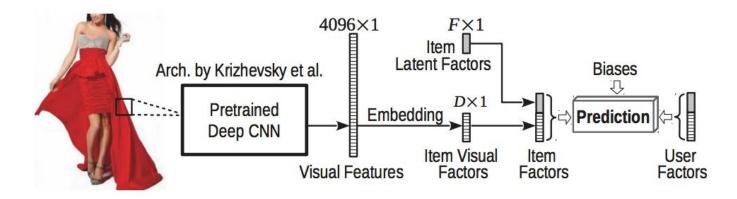


Figure 1: Diagram of our preference predictor. Rating dimensions consist of visual factors and latent (non-visual) factors. Inner products between users and item factors model the compatibility between users and items.

Loss function

$$\hat{x}_{u,i} = \beta_i + \gamma_u^T \gamma_i + \theta_u^T (Ef_i) + \beta^T f_i$$

Weights are learned using BPR-OPT (Rendle et al., 2009)

$$D_S = \{(u, i, j) | u \in U \land i \in I_u^+ \land j \in I \setminus I_u^+ \}$$

$$\sum_{(u,i,j)\in D_S} \ln(\sigma(\hat{x}_{uij}(\Theta))) - \lambda_{\Theta}||\Theta||^2 \qquad \hat{x}_{uij}(\Theta) = \hat{x}_{u,i} - \hat{x}_{u,j}$$

Main Results

AUC	RR	R@20	P@20	nDCG@20	R@100	P@100	nDCG@100
0.75033	0.02351	0.00719	0.00429	0.02717	0.03742	0.00438	0.09059

Example recommendation





















Recommendation (n=20)







































Ground Truth (n=10)





















VBPR Takeaways

Demonstrate usefulness of visual information for implicit feedback recommendation tasks

- Scalable, simple yet effective against cold start

13

Visually-Aware Fashion Recommendation and Design with Generative Image Models

Wang-Cheng Kang UC San Diego wckang@eng.ucsd.edu Chen Fang
Adobe Research
cfang@adobe.com

Zhaowen Wang Adobe Research zhawang@adobe.com Julian McAuley
UC San Diego
jmcauley@eng.ucsd.edu

a.k.a. *DVBPR*:

Deep Visually-aware Bayesian Personalized Ranking

About the paper

- Accepted at the International Conference on Data Mining (ICDM), 2017

- Authors: Adobe Research & Prof. McAuley's Lab @ UCSD

- Proposes fashion 1) recommendation and 2) design (through GANs)

- We focus on 1)

Methodology



Source: Kang et al. 2017

Context

 Fashion domain is complex: long tails, cold starts, evolving dynamics

 Content-aware recommender systems are well-suited to it



Source: Kang et al. 2017

DVBPR Key Insights

- Opt for "domain-aware" visual embeddings instead of "off-the-shelf" as in VBPR

- Joint training of visual embeddings and recommender system

Recommending new items consistent with each user's preference

Approach

- BPR framework: optimize rank of purchased vs non-purchased items

- Siamese trainable CNNs contrast positive and negative pairs
 - Images are retrieved and rescaled in the DataLoader

Original datasets: Amazon fashion + Tradesy

- In this tutorial: updated Wikimedia Commons dataset

Model

- Users $u \in \mathcal{U}$
- Items $i \in \mathcal{I}$
- Positive items \mathcal{I}_n^+
- Item image X_i

VBPR:

$$x_{u,i} = \beta_i + \gamma_u^T \gamma_i + \theta_u^T (E \cdot f_i)$$

DVBPR:

$$x_{u,i} = \beta_i + \gamma_u^T \gamma_i + \theta_u^T \phi(X_i)$$

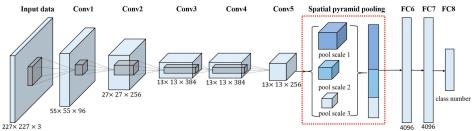
Model

$$x_{u,i} = \beta_i + \gamma_u^T \gamma_i + \theta_u^T \phi(X_i)$$

$$\uparrow \qquad \uparrow$$
We keep this! Trainable CNN

Convolutional Neural Networks

- We use AlexNet with K=100



AlexNet. Source: Han et al. 2017

- Paper uses CNN-F with K=50

conv1	conv2	conv3	conv4	conv5	full6	full7	full8
64x11x11	256x5x5	256x3x3	256x3x3	256x3x3	4096	4096	K
st. 4, pad 0	st. 1, pad 2	st. 1, pad 1	st. 1, pad 1	st. 1, pad 1	drop-	drop-	-
x2 pool	x2 pool	-	-	x2 pool	out	out	-

CNN-F. Source: Kang et al. 2017

BPR Optimization

$$u \in \mathcal{U}$$
 $i \in \mathcal{I}_u^+$ $j \in \mathcal{I} \setminus \mathcal{I}_u^+$ $(u, i, j) \in \mathcal{D}$

$$\mathcal{D} = \{(u, i, j) | u \in \mathcal{U} \land i \in \mathcal{I}_u^+ \land j \in \mathcal{I} \setminus \mathcal{I}_u^+ \}$$

$$\max \sum_{(u,i,j)\in\mathcal{D}} \ln \sigma(x_{u,i,j}) - \lambda_{\Theta} \|\Theta\|^2$$

$$x_{u,i,j} = x_{u,i} - x_{u,j}$$

Retrieval / Recommendation

$$\delta(u, c) = \underset{i \in X_c}{\operatorname{argmax}} \ x_{u,i} = \underset{i \in X_c}{\operatorname{argmax}} \ \beta_i + \gamma_u^T \gamma_i + \theta_u^T \phi(X_i)$$

Observations

- Model converges after 5 epochs (~12 hours on an 8-core CPU + GTX 1080 Ti)

- In our experience, latent CF factors are crucial for the model to learn

Datasets

Dataset	# Users	# Items	# Interactions	# Categories	
Amazon Fashion	64583	234892	513367	6	
Amazon Women	97678	347591	827678	53	
Amazon Men	34244	110636	254870	50	
Tradesy.com	33864	326393	655409	N/A	
Wikimedia (update)	1078	32959	96991	N/A	

- Wikimedia interactions are related to image quality

Main Results

AUC	RR	R@20	P@20	nDCG@20	R@100	P@100	nDCG@100
0.79574	0.07086	0.03164	0.01809	0.09254	0.10099	0.01155	0.14307

Example recommendations

Consumed (n=10)





















Recommendation (n=20)



























































Recommendation (n=20)









































Ground Truth (n=1)



VisRec: A Hands-on Tutorial on Deep Learning for Visual Recommender Systems

DVBPR Takeaways

- Wikimedia dataset is different from Amazon and Tradesy
 - Image quality is primary concern
 - Content < Collaborative Filtering

- This might explain why latent non-visual factors are needed
- DVBPR approach is simple yet effective for visual recommendation in challenging domains
- Newer CNN architectures might be interesting to explore
 - EfficientNet
 - Lambda Networks

Session Conclusions

- BPR is a powerful method to train visual recommender systems

VBPR leverages the power of highly descriptive features that a pre-trained CNN can generate

- DVBPR takes this one step further by *learning* the weights of the CNN module

References (VBPR)

[1] He, R., & McAuley, J. (2016, February). VBPR: visual bayesian personalized ranking from implicit feedback. In Proceedings of the AAAI Conference on Artificial Intelligence (Vol. 30, No. 1).

[2] Messina, P., Dominguez, V., Parra, D., Trattner, C., & Soto, A. (2019). Content-based artwork recommendation: integrating painting metadata with neural and manually-engineered visual features. User Modeling and User-Adapted Interaction, 29(2), 251-290.

[3] Rendle, S., Freudenthaler, C., Gantner, Z., & Schmidt-Thieme, L. (2012). BPR: Bayesian personalized ranking from implicit feedback. arXiv preprint arXiv:1205.2618.

References (DVBPR)

[1] Kang, W., Fang, C., Wang, Z., & McAuley, J. (2017). Visually-Aware Fashion Recommendation and Design with Generative Image Models. *2017 IEEE International Conference on Data Mining (ICDM)*, 207-216.

[2] Krizhevsky, A., Sutskever, I., & Hinton, G.E. (2012). ImageNet classification with deep convolutional neural networks. *Communications of the ACM, 60*, 84 - 90.

[3] Chatfield, K., Simonyan, K., Vedaldi, A., & Zisserman, A. (2014). Return of the Devil in the Details: Delving Deep into Convolutional Nets. *ArXiv*, abs/1405.3531.











Session 3: Visual BPR methods

Questions? Shoot me an email at pcerdam@uc.cl

Denis Parra¹, Antonio Ossa-Guerra¹, Manuel Cartagena¹, **Patricio Cerda-Mardini**², Felipe del Río¹, Isidora Palma¹, Diego Saez-Trumper³, and Miriam Redi³

- 1. Pontificia Universidad Católica de Chile
- 2. mindsdb
- 3. Wikimedia Foundation



Appendix - Implementation details

Implementation details

```
VisualRecSys-Tutorial-IUI2021 / models / vbpr.py / > Jump to -
    aaossa 🖖 Add VBPR model implementation ....
८३ 1 contributor
123 lines (98 sloc) | 4.51 KB
      """VBPR implementation in PyTorch
      import torch
      import torch.nn as nn
      import torch.nn.functional as F
      class VBPR(nn.Module):
          """VBPR model architecture from 'VBPR: Visual Bayesian
          Personalized Ranking from Implicit Feedback'.
```

Implementation details (Bias terms)

```
def forward(self, ui, pi, ni):
class VBPR(nn.Module):
   """VBPR model architecture from 'VBPR: Visual Bayesian
   Personalized Ranking from Implicit Feedback'.
                                                                                                         ui latent factors = self.gamma users(ui) # Latent factors of user u
                                                                                                         ui_visual_factors = self.theta_users(ui) # Visual factors of user u
   def init (self, n users, n items, features, dim gamma, dim thet
                                                                                                         pi_bias = self.beta_items(pi) # Pos. item bias
       super(). init ()
                                                                                                         ni bias = self.beta items(ni) # Neg. item bias
                                                                                                         pi_latent_tactors = self.gamma_items(pi) # Pos. item visual factors
                                                                                                         ni latent factors = self.gamma items(ni) # Neg. item visual factors
       self.features = nn.Embedding.from pretrained(features, freeze=Tru
                                                                                                         pi_features = self.features(pi) # Pos. item visual features
                                                                                                         ni features = self.features(ni) # Neg. item visual features
       self.gamma users = nn.Embedding(n users, dim gamma)
       self.gamma items = nn.Embedding(n items, dim gamma)
                                                                                                         diff_features = pi_features - ni_features
                                                                                                         diff latent factors = pi latent factors - ni latent factors
       self.theta users = nn.Embedding(n users, dim theta)
       self.embedding = nn.Embedding(features.size(1), dim_theta)
                                                                                                             pi bias - ni bias
       # Biases (beta)
                                                                                                             + (ui latent factors * diff latent factors).sum(dim=1).unsqueeze(-1)
                                                                                                             + (ui visual factors * diff features.mm(self.embedding.weight)).sum(dim
       self.beta items = nn.Embedding(n items, 1)
                                                                                                             + diff features.mm(self.visual bias.weight)
       self.visual bias = nn.Embedding(features.size(1), 1)
       self.reset_parameters()
```

Implementation details (Latent factors)

```
def forward(self, ui, pi, ni):
class VBPR(nn.Module):
   """VBPR model architecture from 'VBPR: Visual Bayesian
   Personalized Ranking from Implicit Feedback'.
                                                                                                         ui latent factors = self.gamma users(ui) # Latent factors of user u
                                                                                                         ui_visual_factors = self.theta_users(ui) # Visual_factors of user u
   def init (self, n users, n items, features, dim gamma, dim theta)
                                                                                                         pi_bias = self.beta_items(pi) # Pos. item bias
       super(). init ()
                                                                                                         ni bias = self.beta items(ni) # Neg. item bias
                                                                                                         pi_latent_factors = self.gamma_items(pi) # Pos. item visual factors
                                                                                                         ni latent factors = self.gamma items(ni) # Neg. item visual factors
       self.features = nn.Embedding.from pretrained(features, freeze=Tru
                                                                                                         pi_features = self.features(pi) # Pos. item visual features
                                                                                                         ni features = self.features(ni) # Neg. item visual features
       self.gamma users = nn.Embedding(n users, dim gamma)
        self.gamma items = nn.Embedding(n items, dim gamma)
                                                                                                         diff features = pi features - ni features
                                                                                                         diff latent factors = pi latent factors - ni latent factors
       self.theta users = nn.Embedding(n users, dim theta)
       self.embedding = nn.Embedding(features.size(1), dim_theta)
                                                                                                            pi bias - ni bias
       # Biases (beta)
                                                                                                             + (ui latent factors * diff latent factors).sum(dim=1).unsqueeze(-1)
                                                                                                             + (ui visual factors * diff features.mm(self.embedding.weight)).sum(dim
       self.beta items = nn.Embedding(n items, 1)
                                                                                                             + diff features.mm(self.visual bias.weight)
       self.visual bias = nn.Embedding(features.size(1), 1)
       # Random weight initialization
       self.reset_parameters()
                                                                                                         return x uij.unsqueeze(-1)
```

Implementation details (Visual factors)

```
def forward(self, ui, pi, ni):
class VBPR(nn.Module):
   """VBPR model architecture from 'VBPR: Visual Bayesian
                                                                                                         ui latent factors = self.gamma users(ui) # Latent factors of user u
   Personalized Ranking from Implicit Feedback'.
                                                                                                         ui_visual_factors = self.theta_users(ui) # Visual factors of user u
   def init (self, n users, n items, features, dim gamma, dim the
                                                                                                         pi_bias = self.beta_items(pi) # Pos. item bias
       super(). init ()
                                                                                                         ni bias = self.beta items(ni) # Neg. item bias
                                                                                                         pi_latent_factors = self.gamma_items(pi) # Pos. item visual factors
                                                                                                         ni latent factors = self.gamma items(ni) # Neg. item visual factors
       self.features = nn.Embedding.from pretrained(features, freeze=1ru
                                                                                                         pi_features = self.features(pi) # Pos. item visual features
                                                                                                         ni features = self.features(ni) # Neg. item visual features
       self.gamma users = nn.Embedding(n users, dim gamma)
       self.gamma items = nn.Embedding(n items, dim gamma)
                                                                                                         diff_features = pi_features - ni_features
                                                                                                         diff latent factors = pi latent factors - ni latent factors
       self.theta users = nn.Embedding(n users, dim theta)
       self.embedding = nn.Embedding(features.size(1), dim_theta)
                                                                                                         x uij = (
                                                                                                             pi bias - ni bias
       # Biases (beta)
                                                                                                             + (ui latent factors * diff latent factors).sum(dim=1).unsqueeze(-1)
                                                                                                             + (ui visual factors * diff features.mm(self.embedding.weight)).:um(di
       self.beta items = nn.Embedding(n items, 1)
                                                                                                             + diff features.mm(self.visual bias.weight)
       self.visual bias = nn.Embedding(features.size(1), 1)
       # Random weight initialization
       self.reset_parameters()
                                                                                                         return x uij.unsqueeze(-1)
```

Implementation details (Users' visual bias)

```
def forward(self, ui, pi, ni):
class VBPR(nn.Module):
   """VBPR model architecture from 'VBPR: Visual Bayesian
   Personalized Ranking from Implicit Feedback'.
                                                                                                         ui latent factors = self.gamma users(ui) # Latent factors of user u
                                                                                                         ui_visual_factors = self.theta_users(ui) # Visual factors of user u
   def init (self, n users, n items, features, dim gamma, dim theta)
                                                                                                         pi_bias = self.beta_items(pi) # Pos. item bias
       super(). init ()
                                                                                                         ni bias = self.beta items(ni) # Neg. item bias
                                                                                                         pi_latent_factors = self.gamma_items(pi) # Pos. item visual factors
                                                                                                         ni latent factors = self.gamma items(ni) # Neg. item visual factors
       self.features = nn.Embedding.from pretrained(features, freeze=Tru
                                                                                                         pi features = self.features(pi) # Pos. item visual features
                                                                                                         ni features = self.features(ni) # Neg. item visual features
       self.gamma users = nn.Embedding(n users, dim gamma)
       self.gamma items = nn.Embedding(n items, dim gamma)
                                                                                                         diff features = pi features - ni features
                                                                                                         diff latent factors = pi latent factors - ni latent factors
       self.theta users = nn.Embedding(n users, dim theta)
       self.embedding = nn.Embedding(features.size(1), dim theta)
                                                                                                             pi bias - ni bias
                                                                                                             + (ui latent factors * diff latent factors).sum(dim=1).unsqueeze(-1)
                                                                                                             + (ui visual factors * diff features.mm(self.embedding.weight)).sum(dim
       self.beta items = nn.Fmbedding(n items. 1)
                                                                                                             + diff features.mm(self.visual bias.weight)
       self.visual bias = nn.Embedding(features.size(1), 1)
       # Random weight initialization
                                                                                                         return x uij.unsqueeze(-1)
       self.reset_parameters()
```

Implementation details (Inference)

At this time, we don't need the first 2 terms, because they're constants shared across the recommendation list

$$\widehat{x}_{u,i} = \alpha + \beta_u + \beta_i + \gamma_u^T \gamma_i + \theta_u^T (\mathbf{E}f_i) + \beta'^T f_i.$$

```
def recommend all(self, user, cache=None, grad enabled=False):
    with torch.set grad enabled(grad enabled):
       u latent factors = self.gamma users(user) # Latent factors of user u
       u visual factors = self.theta users(user) # Visual factors of user u
       i bias = self.beta items.weight # Items bias
        i latent factors = self.gamma items.weight # Items visual factors
       i_features = self.features.weight # Items visual features
       if cache is not None:
           visual rating space, opinion visual appearance = cache
           visual rating space = i features.mm(self.embedding.weight)
           opinion visual appearance = i features.mm(self.visual bias.weight)
            + (u_latent_factors * i_latent_factors).sum(dim=1).unsqueeze(-1)
            + (u visual factors * visual rating space).sum(dim=1).unsqueeze(-1)
           + opinion visual appearance
       return x ui
```

Implementation details - Model Class

```
class DVBPR(nn.Module):
   def init (self, n users, n items, K=2048, use cnnf=False):
       super(). init ()
       self.cache = None
       # CNN for learned image features
       if use cnnf:
            self.cnn = CNNF(hidden dim=K) # CNN-F is a smaller CNN
       else:
           alexnet = models.alexnet(pretrained=False)
            final len = alexnet.classifier[-1].weight.shape[1]
            alexnet.classifier[-1] = nn.Linear(final len, K)
            self.cnn = alexnet
       # Visual latent preference (theta)
        self.theta users = nn.Embedding(n users, K)
       # Latent factors (gamma)
        self.gamma users = nn.Embedding(n users, 100)
        self.gamma items = nn.Embedding(n items, 100)
       # Random weight initialization
        self.reset parameters()
```

Implementation details - Forward Pass

```
def forward(self, ui, pimg, nimg, pi, ni):
    ui visual factors = self.theta users(ui) # Visual factors of user u
    ui latent factors = self.gamma users(ui) # Latent factors of user u
    # Items
    pi features = self.cnn(pimg) # Pos. item visual features
    ni features = self.cnn(nimg) # Neg. item visual features
    pi latent factors = self.gamma items(pi) # Pos. item visual factors
    ni latent factors = self.gamma items(ni) # Neg. item visual factors
    x ui = (ui visual factors * pi features).sum(1) + (pi latent factors * ui latent factors).sum(1)
    x uj = (ui visual factors * ni features).sum(1) + (ni latent factors * ui latent factors).sum(1)
    return x ui, x uj
```

Implementation details - Optimization

```
# Forward pass
with torch.set_grad_enabled(phase == "train"):
    pos, neg = self.model(profile, pimg, nimg, pi, ni)
    output = pos-neg
    loss = self.criterion(output, target)
    loss += (1.0 * torch.norm(self.model.theta_users.weight))

# Backward pass
if phase == "train":
    loss.backward()
    self.optimizer.step()
```

Implementation details - Recommendation

```
recommend all(self, user, cache, grad enabled=False):
with torch.set grad enabled(grad enabled):
   # User
    u visual factors = self.theta users(user) # Visual factors of user u
    ui latent factors = self.gamma users(user)
   # Ttems
    i latent factors = self.gamma items.weight # Items visual factors
    x ui = ((i latent factors * ui latent factors).sum(dim=1).squeeze() + \
            (u visual factors * cache).sum(dim=2).squeeze())
    return x ui
```